

# **A taste of Big Data**

Raúl Ramos Pollán

Universidad Industrial de Santander

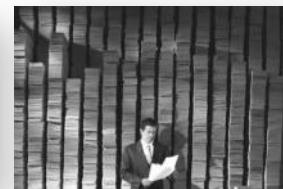
# Cómputo escalable

## Algoritmos/Aplicaciones

```
Initialize Q(s, a) arbitrarily
Repeat (for each episode):
    Initialize s
    Choose a from s using policy derived from Q
        (e.g., ε-greedy)
    Repeat (for each step of episode):
        Take action a, observe r, s'
        Choose a' from s' using policy derived from Q
            (e.g., ε-greedy)
        Q(s, a) ← Q(s, a) + α[r + γQ(s', a') - Q(s, a)]
        s ← s'; a ← a';
    until s is terminal
```



## Datos



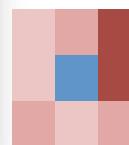
## Infraestructura de cómputo



## Recurso humano/ Comunidades científicas

# Paralelización

Procesamiento de imágenes → operaciones en regiones (i.e. contraste)

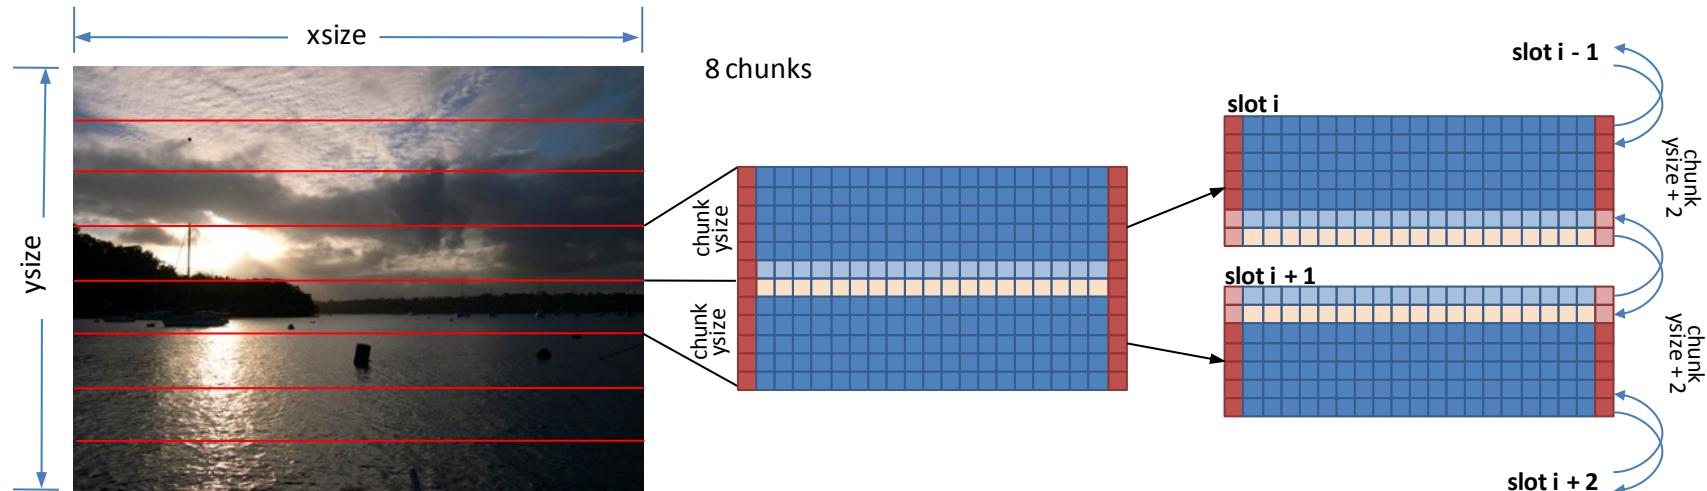


and perform n iterations

# Ejemplos de algoritmos

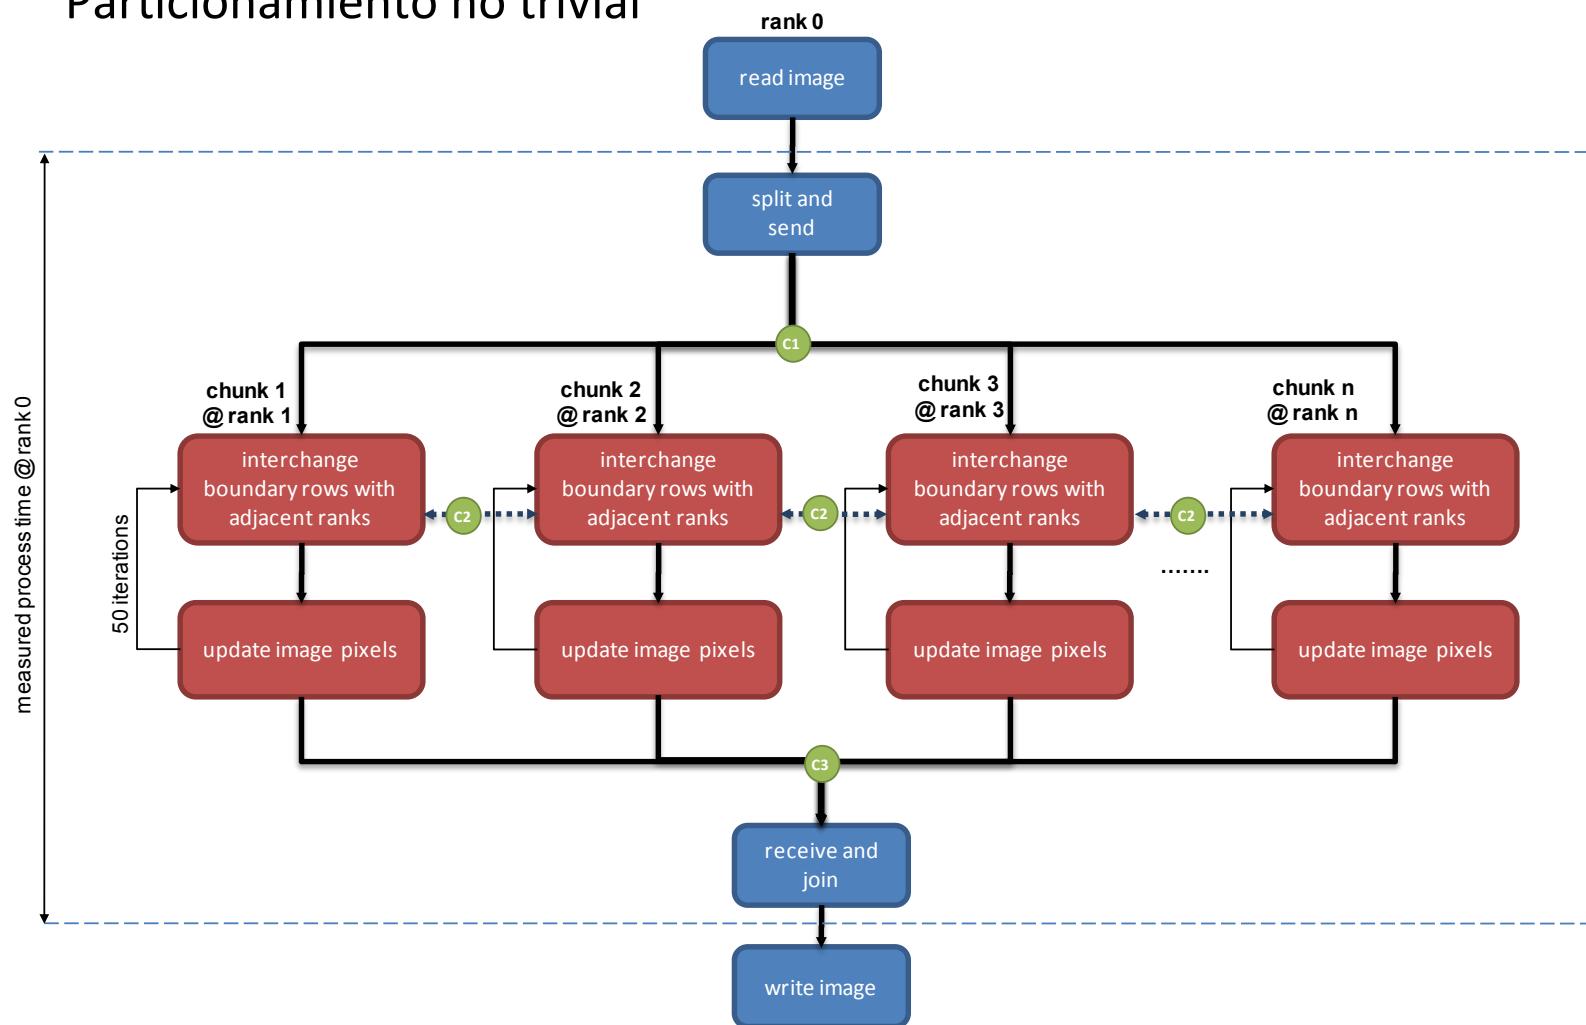
Particionamiento no trivial → por algoritmo

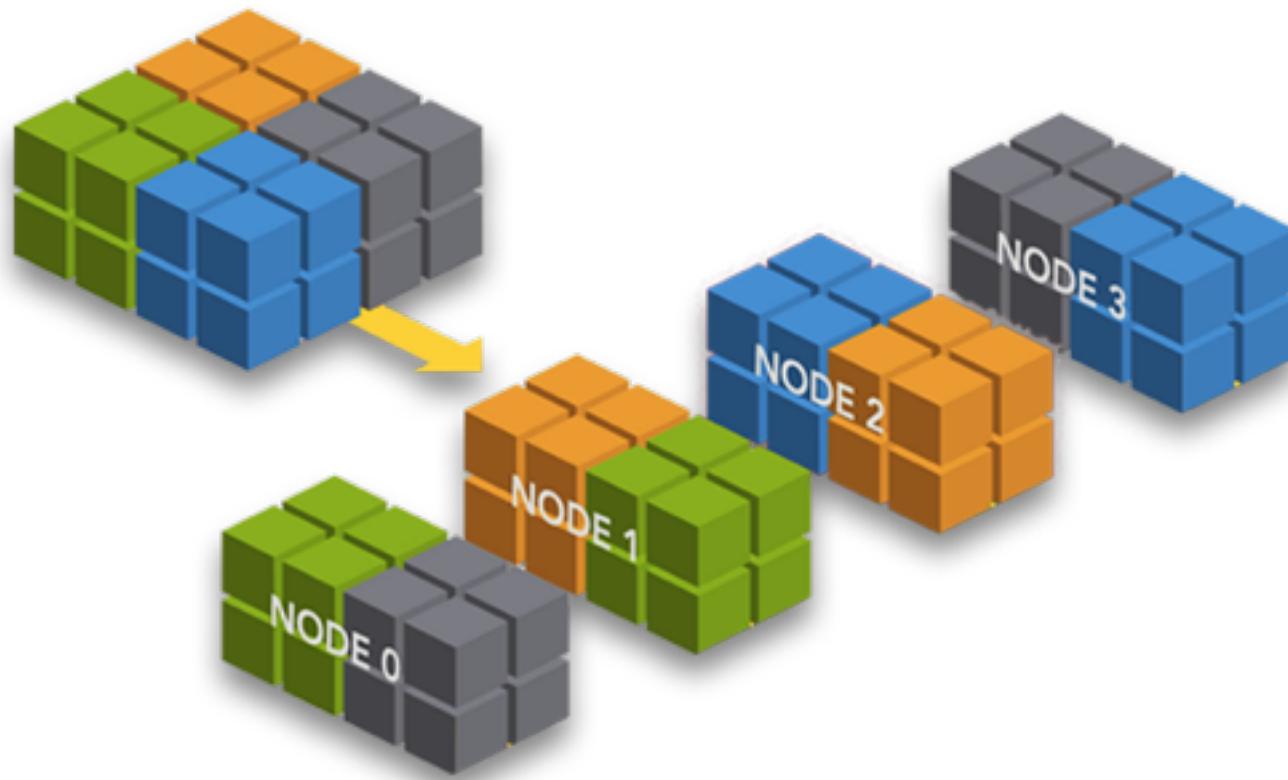
(otra estrategia → por datos para muchas imágenes, pero cada una tiene que caber en memoria)



# Ejemplos de algoritmos

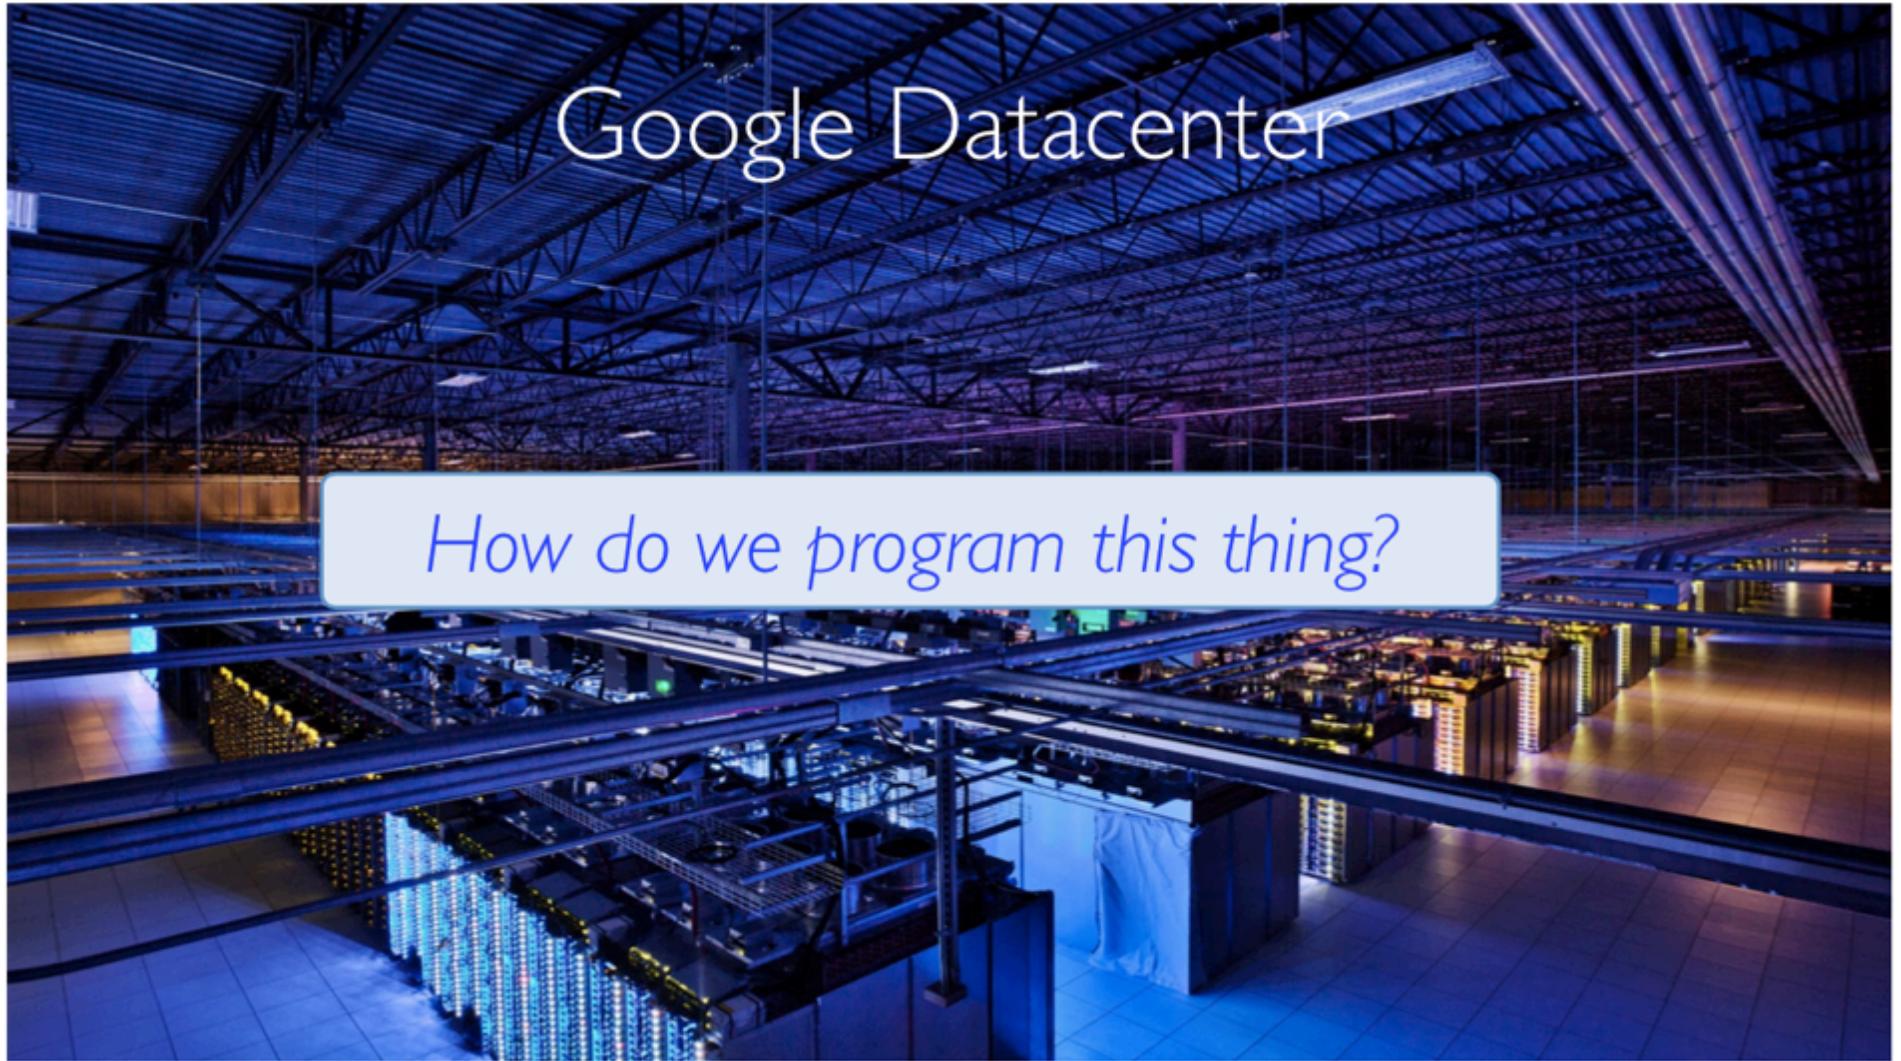
## Particionamiento no trivial



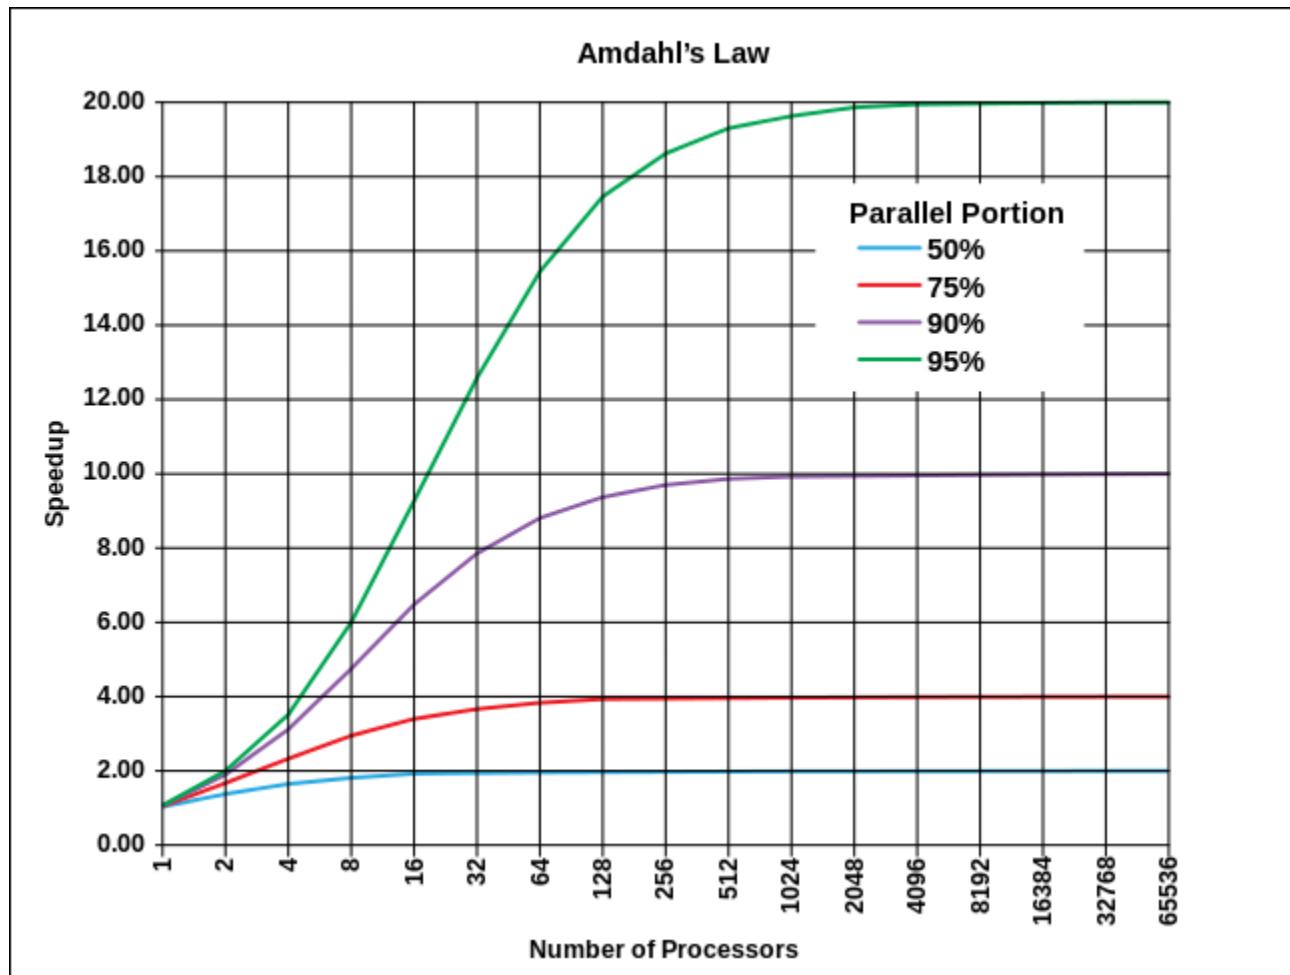


# Google Datacenter

*How do we program this thing?*



# Amdahl's law



# Parallelize computation of the mean

$$X = \{x_i\} \quad i \in [1, n]$$

$$\bar{x} = \text{mean}(X) = \frac{1}{m} \sum_{i=1}^n x_i$$

# Parallelize computation of the mean

$$X = \{x_i\} \quad i \in [1, n]$$

$$\bar{x} = \text{mean}(X) = \frac{1}{m} \sum_{i=1}^n x_i$$

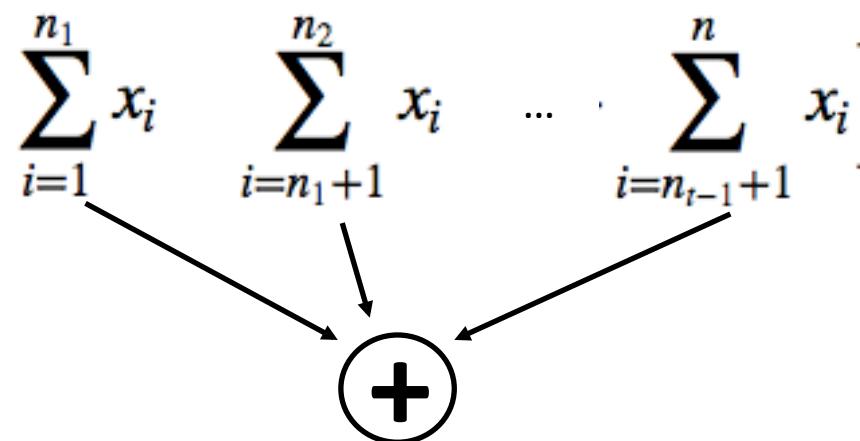
$$\bar{x} = \frac{1}{m} \left[ \sum_{i=1}^{n_1} x_i + \sum_{i=n_1+1}^{n_2} x_i + \dots + \sum_{i=n_{t-1}+1}^n x_i \right]$$

# Parallelize computation of the mean

$$X = \{x_i\} \quad i \in [1, n]$$

$$\bar{x} = \text{mean}(X) = \frac{1}{m} \sum_{i=1}^n x_i$$

$$\bar{x} = \frac{1}{m} \left[ \sum_{i=1}^{n_1} x_i + \sum_{i=n_1+1}^{n_2} x_i + \dots + \sum_{i=n_{t-1}+1}^n x_i \right]$$



## How to scale?

$$X = \{x_i\} \quad i \in [1, n]$$

$$\bar{x} = \text{mean}(X) = \frac{1}{m} \sum_{i=1}^n x_i \quad \text{var}(X) = \frac{1}{m} \sum_{i=1}^n (\bar{x} - x_i)^2$$

## How to scale?

$$X = \{x_i\} \quad i \in [1, n]$$

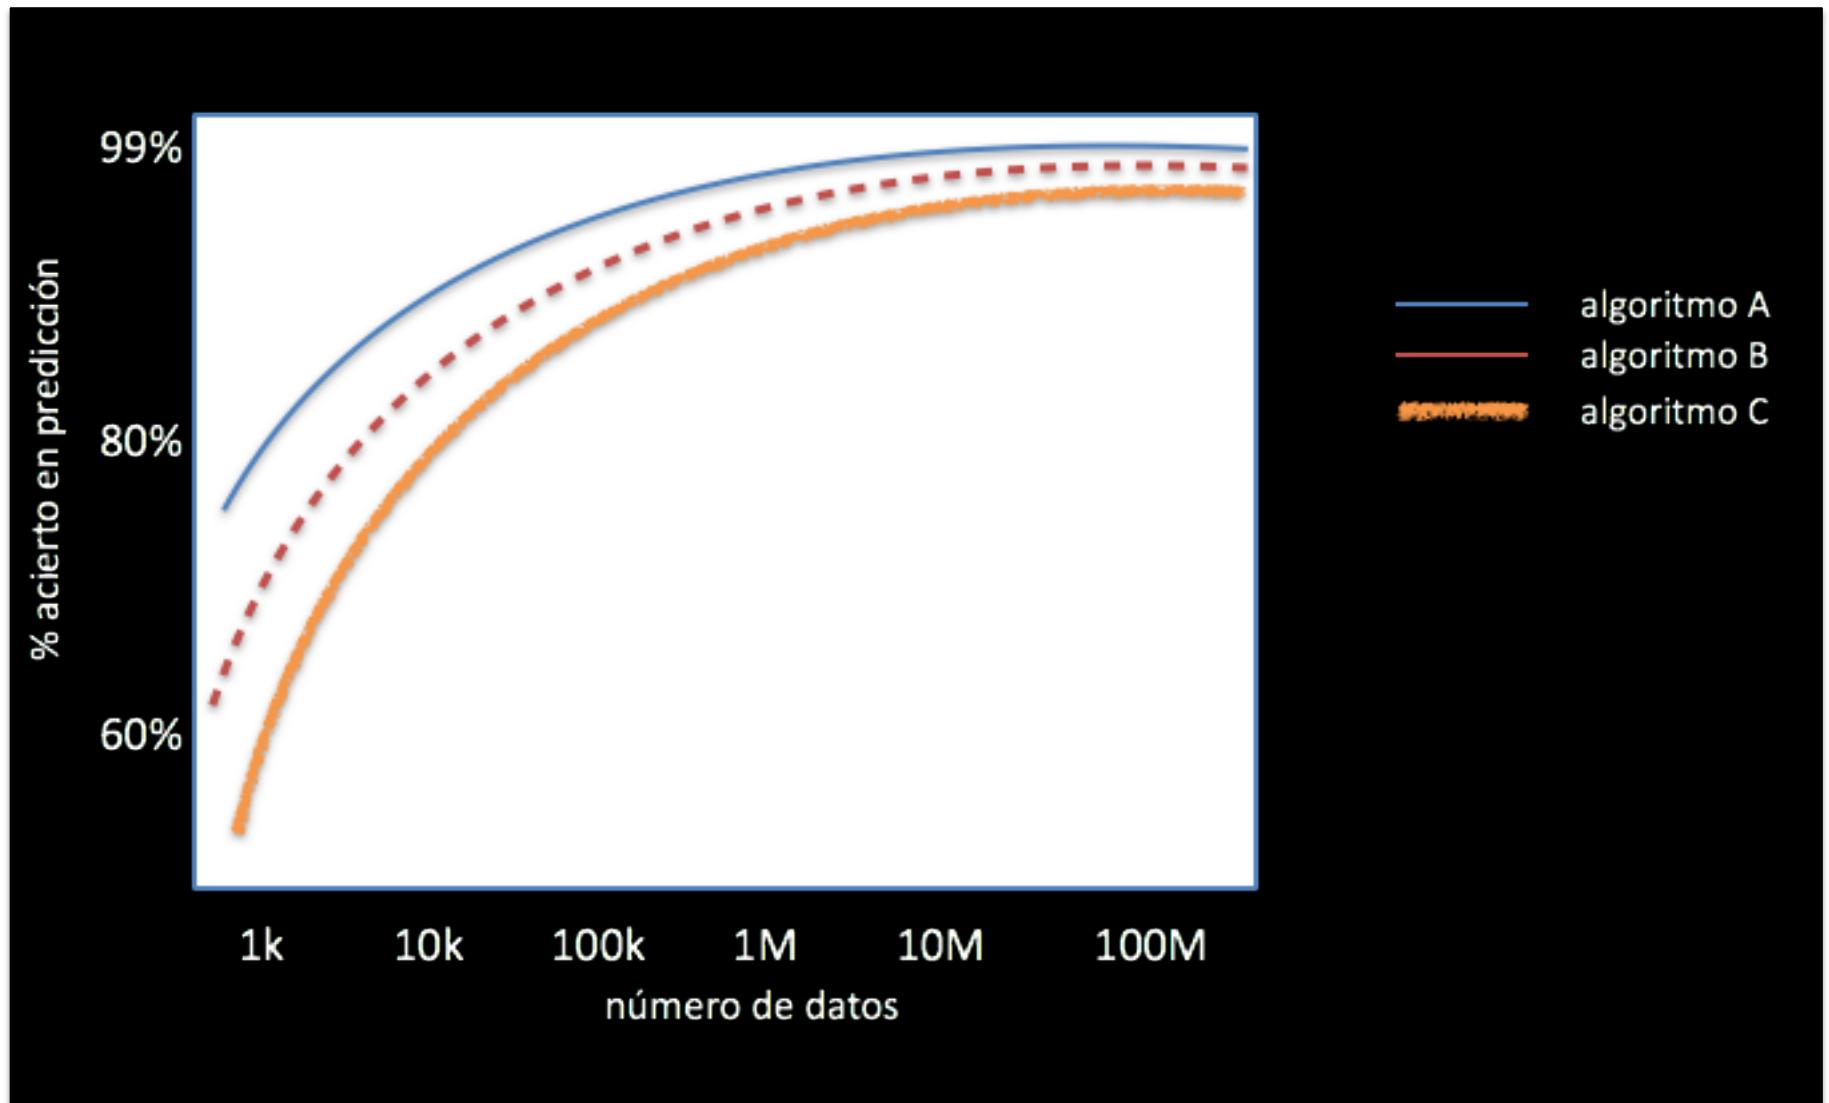
$$\bar{x} = mean(X) = \frac{1}{m} \sum_{i=1}^n x_i \quad var(X) = \frac{1}{m} \sum_{i=1}^n (\bar{x} - x_i)^2$$

$$var(X) = \frac{1}{m} \left[ \sum \bar{x}^2 - 2\bar{x} \sum x_i + \sum x_i^2 \right]$$

$$var(X) = \frac{1}{m} \left[ \frac{1}{m} \left( \sum x_i \right)^2 - \frac{1}{m} 2 \left( \sum x_i \right)^2 + \sum x_i^2 \right]$$

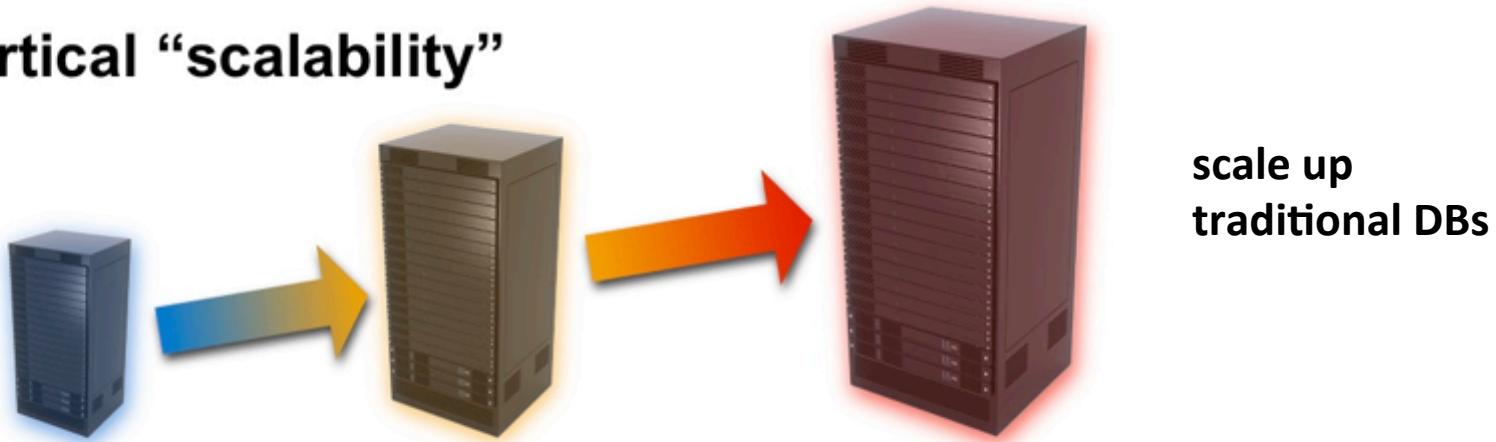
$$var(X) = \frac{1}{m} \sum x_i^2 - \frac{1}{m^2} \left( \sum x_i \right)^2$$

# Predictive Analytics in Big Data



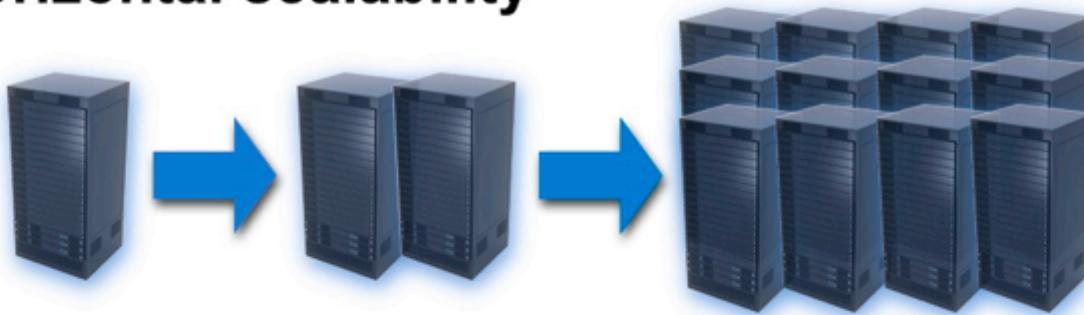
# Scalability in Big Data

## Vertical “scalability”



scale up  
traditional DBs

## Horizontal scalability



scale out  
noSQL

seek “triviality” for appropriate  
sw+hw architectures

recent technologies (virtualization, etc.) tend to favor the cost of scaling out

## **How to scale?**

**Parallelize the code yourself!!**



**Lots of flexibility .... but HARD!!!**

# map reduce

```
2012-01-01 09:08 BOG Libros 88.56 Discover
2012-01-01 09:09 BGA Libros 337.71 Efectivo
2012-01-01 09:52 BGA Libros 62.41 Discover
2012-01-01 10:08 MED Musica 93.37 Visa
2012-01-01 10:22 BGA Musica 369.94 MasterCard
2012-01-01 10:58 MED Musica 119.12 Efectivo
2012-01-01 11:36 BOG Musica 296.76 Discover
```

# map reduce

BOG 88.56

BGA 337.71

BGA 62.41

MED 93.37

BGA 369.94

MED 119.12

BOG 296.76

BOG 88.56 296.76

BGA 337.71 62.41 369.94

MED 93.37 119.12

BOG 385.32

BGA 770.06

MED 212.49

# map reduce

```
map (input record):  
    yield key=record[2], value=record[4]
```

shuffle → gathers all values for each key

```
reduce (key, [values]):  
    yield key, sum(values)
```

# map reduce

```
map (input record):  
    yield key1, value1
```

shuffle

```
reduce (key1, [list of values]):  
    yield key2, value2
```



## How to crunch 1PB?

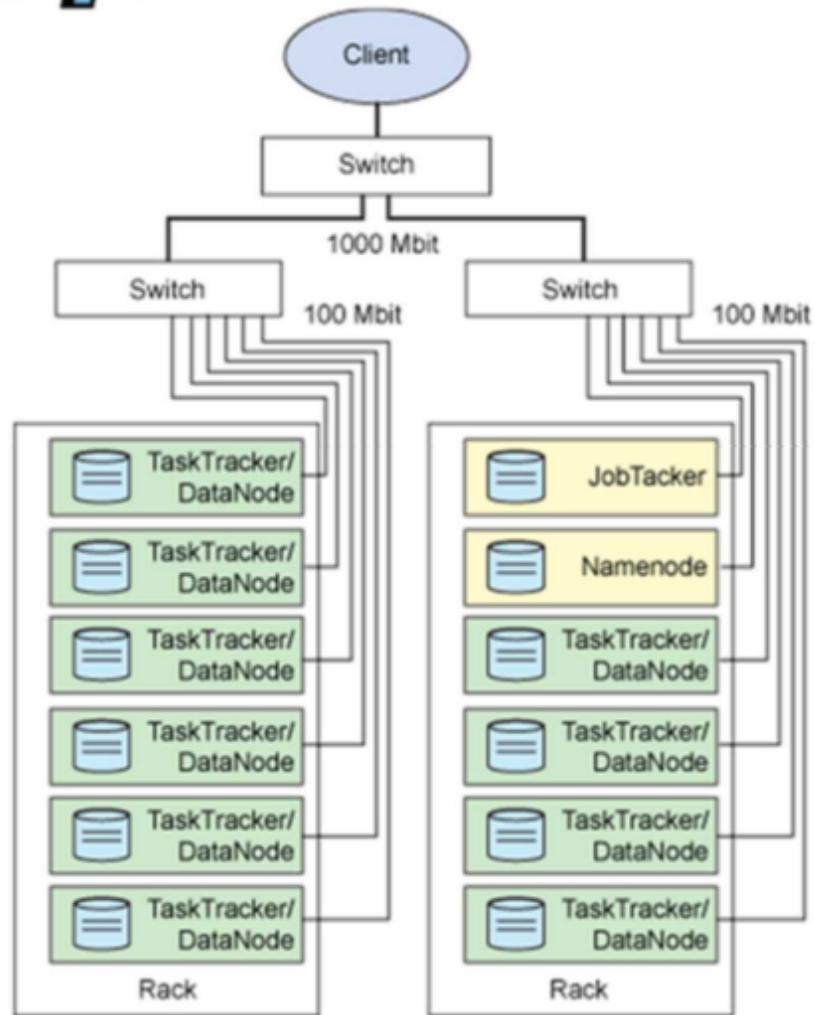
Lots of disks spinning all the time  
Redundancy, since disks die  
Lots of CPU cores, working all the time  
Retry, since network errors happen

## Design Qualities

Scalable – many servers with cores and disks  
Reliable – redundant storage  
Fault-tolerant – auto retry, self healing

## Computation to Data

Very simple computing model □ Map-Reduce  
Each computing node is also a storage node  
HDFS □ on top of ext3, fixed 64MB file blocks  
write once, read many



**Data goes to computing**



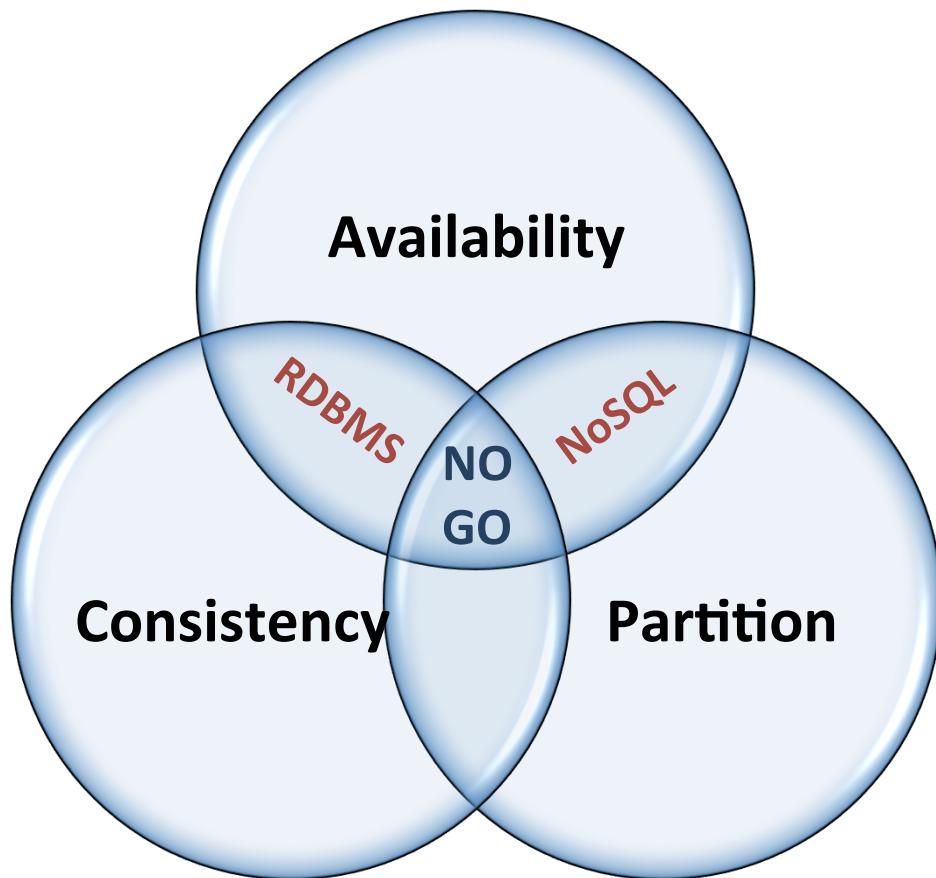
**Computing goes to data**

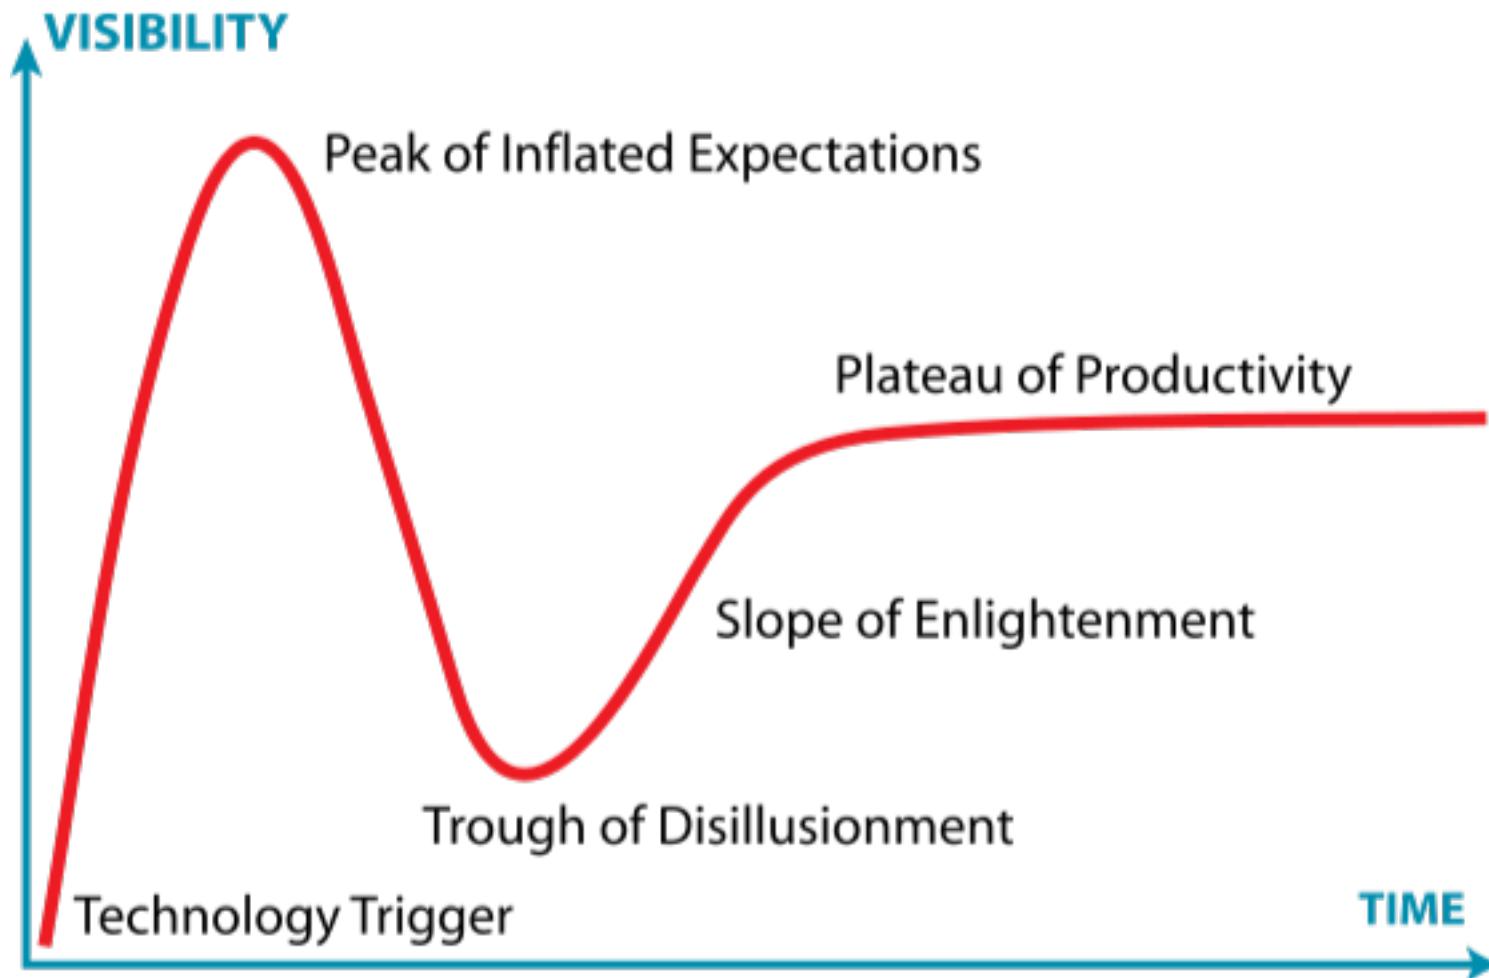
# NoSQL

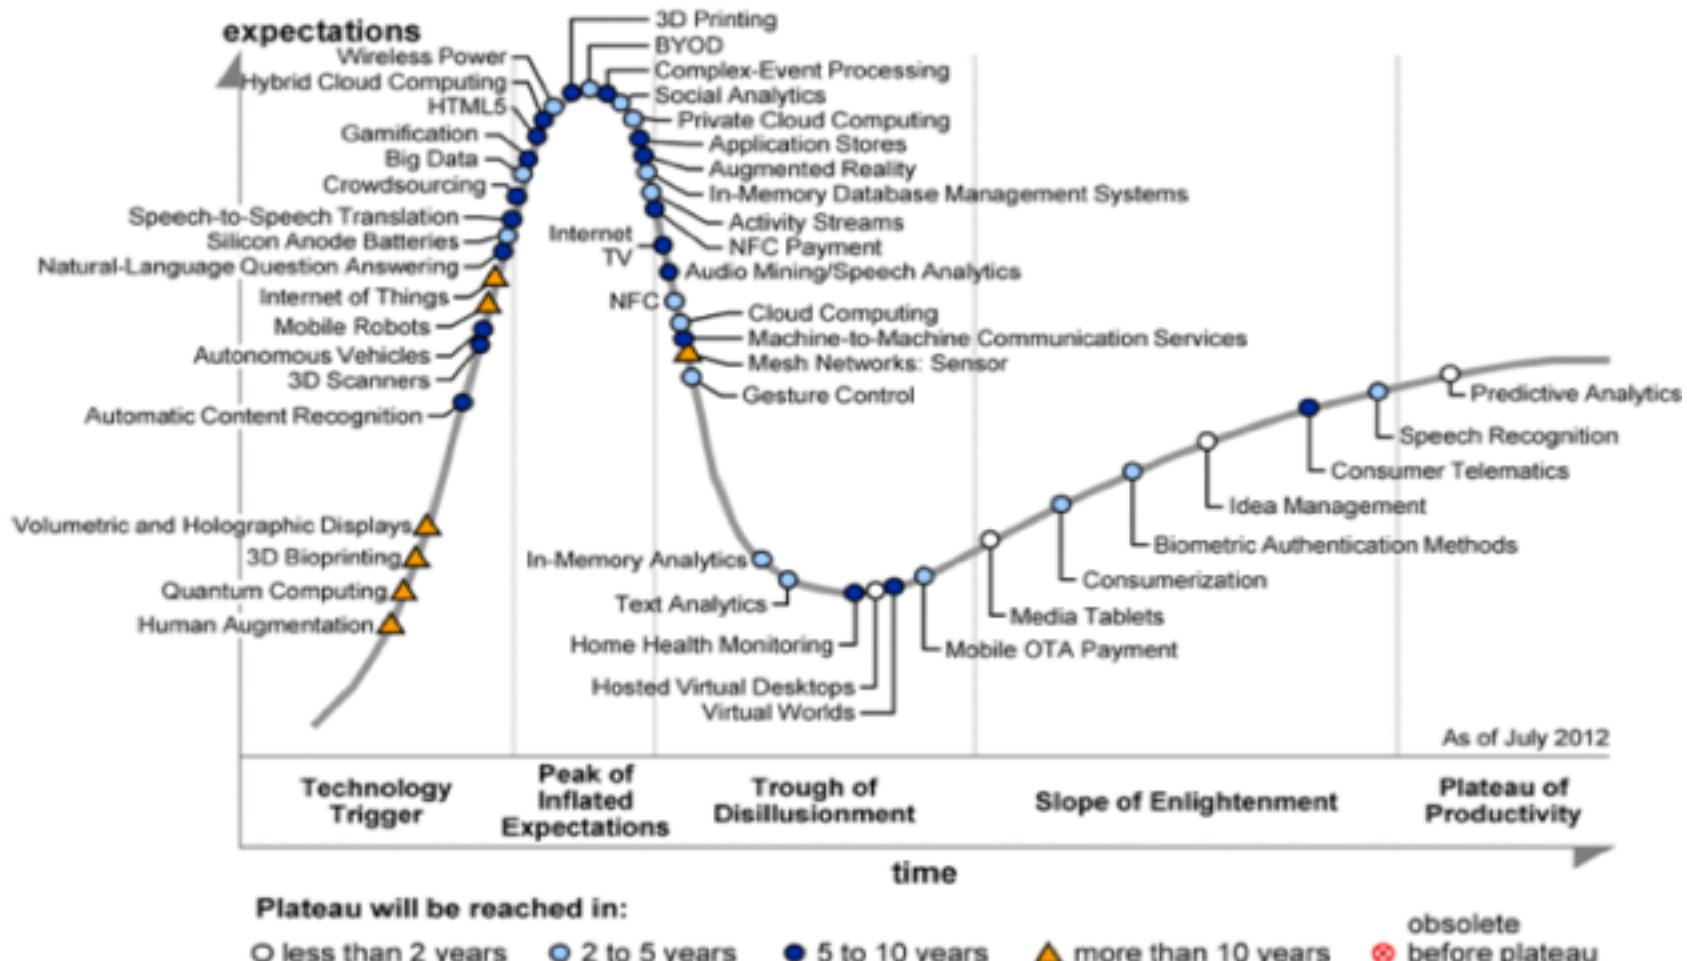
- Expressivity SQL vs. Scalability
- Simpler data model (key, values)
- Simpler operations
  - Scan/access per key, basic transactions (check&put)  
No joins, no SQL language
- Simple failover and scale up
- Big table, Hbase, DynamoDB, Azure, Cassandra, etc.

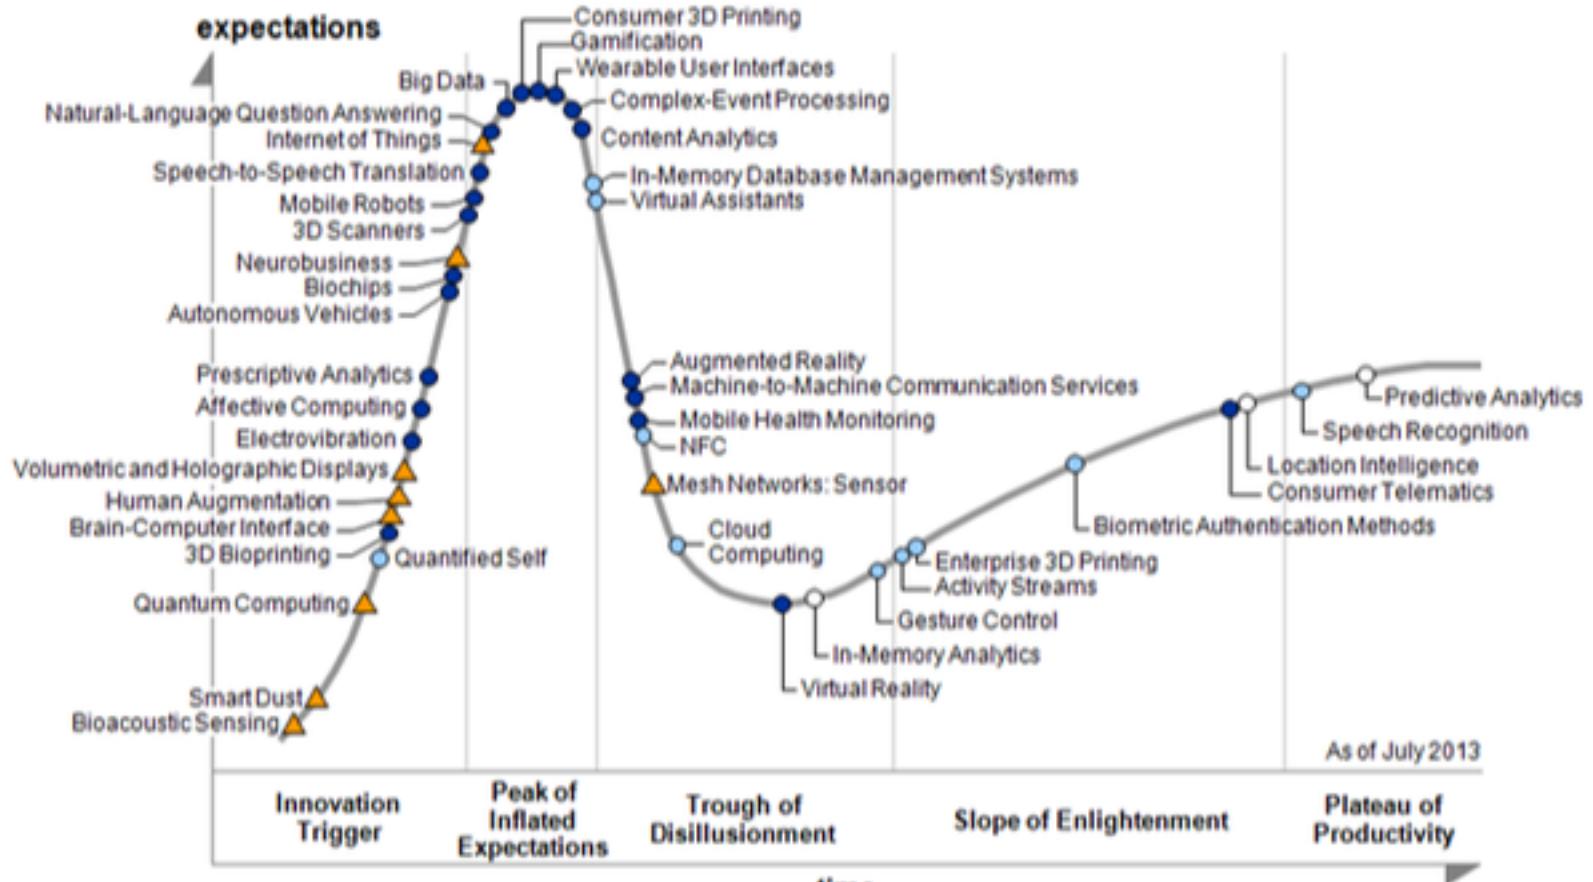
→ more work for programmers!!!

# NoSQL Eventual Consistency









Plateau will be reached in:

- less than 2 years
- 2 to 5 years
- 5 to 10 years
- ▲ more than 10 years
- ✖ obsolete
- ✖ before plateau

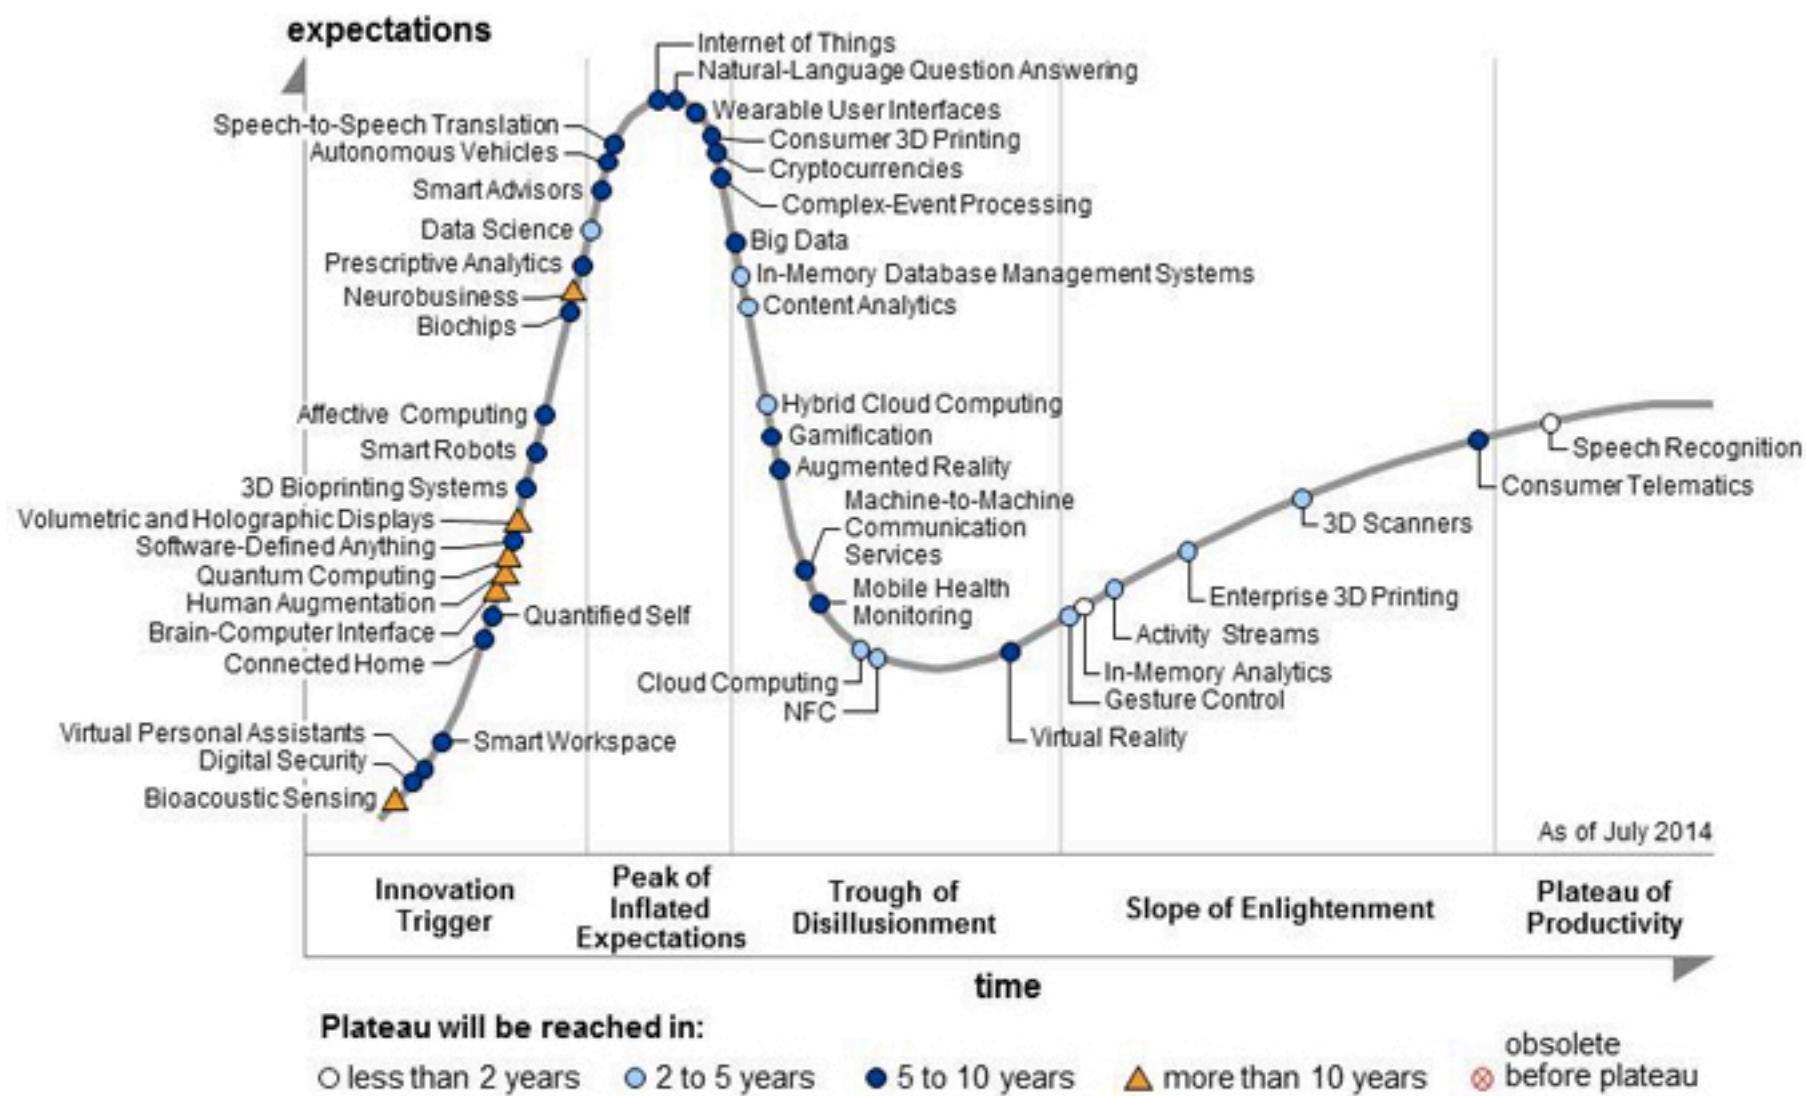
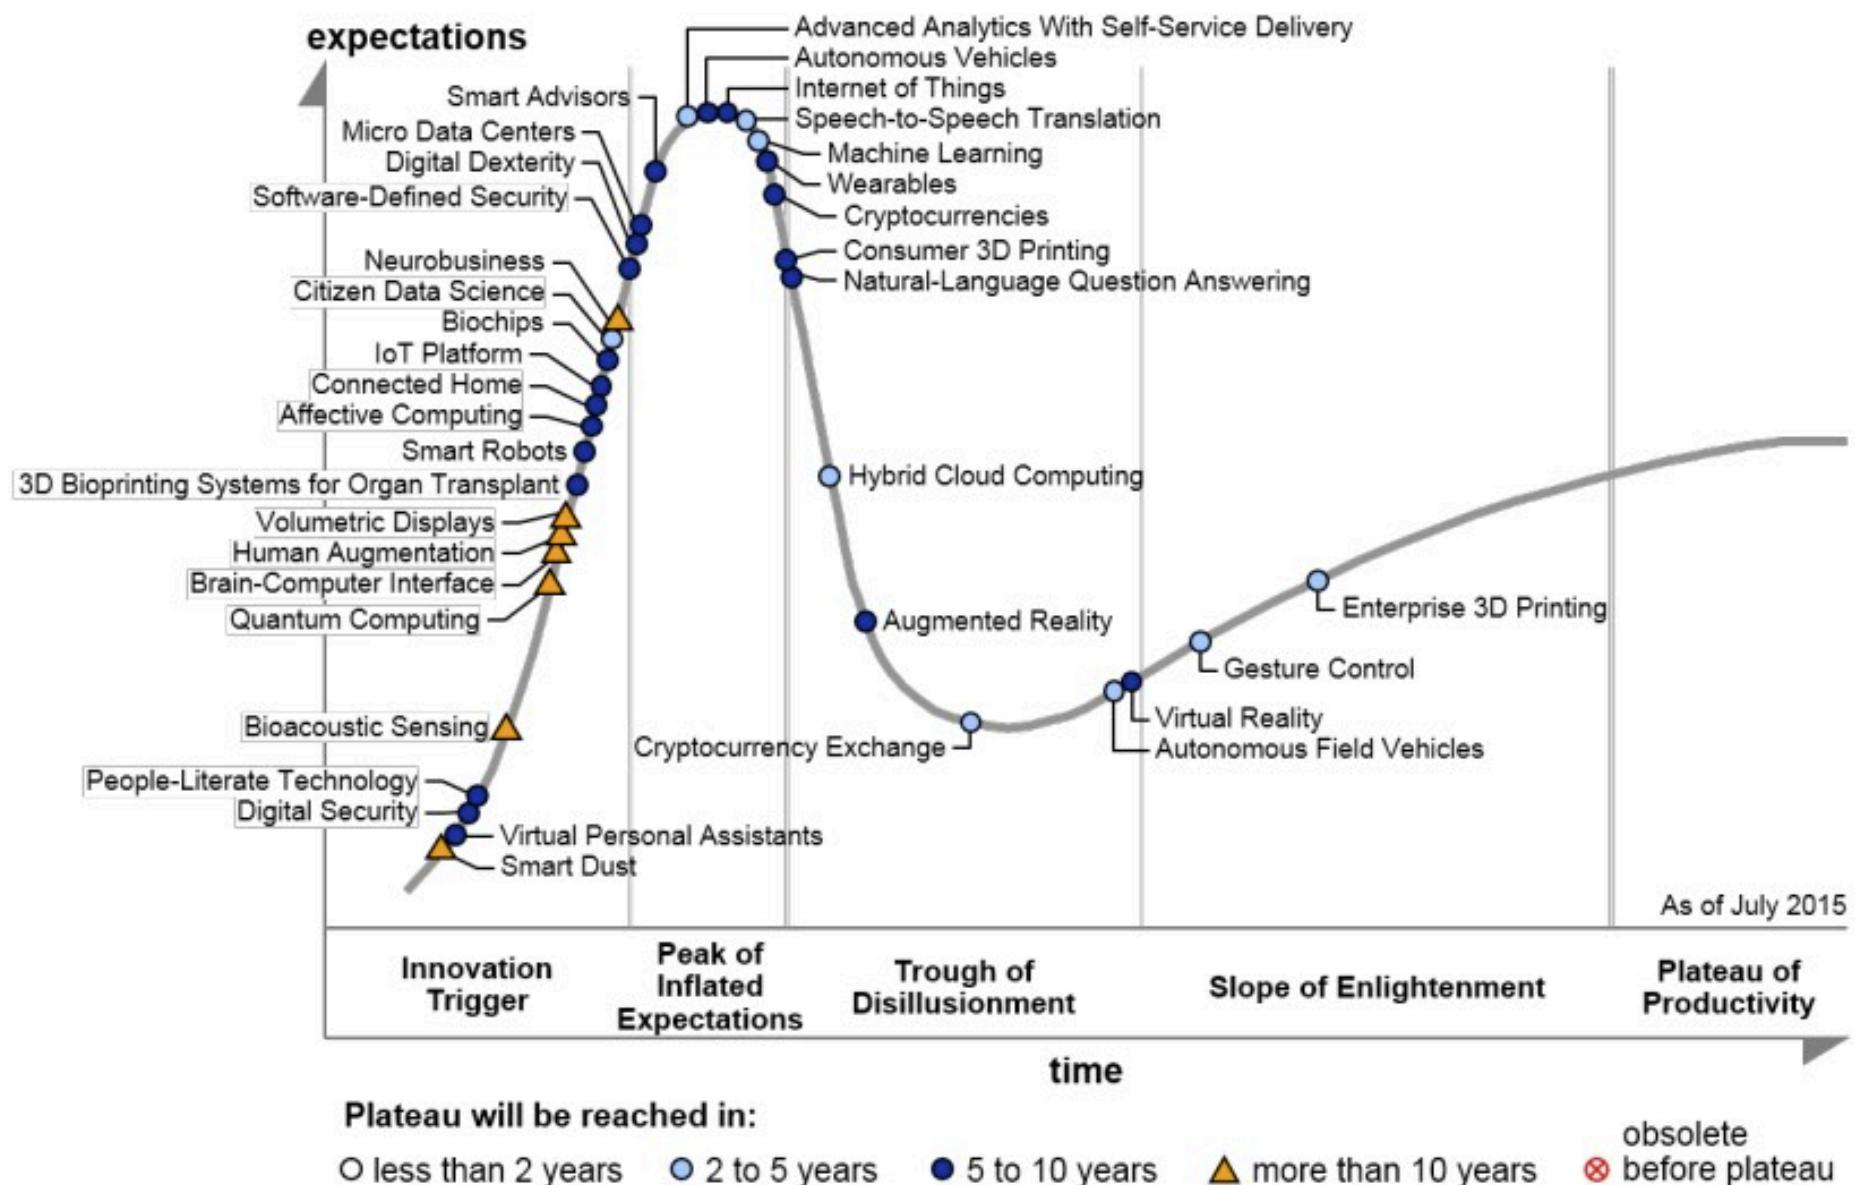


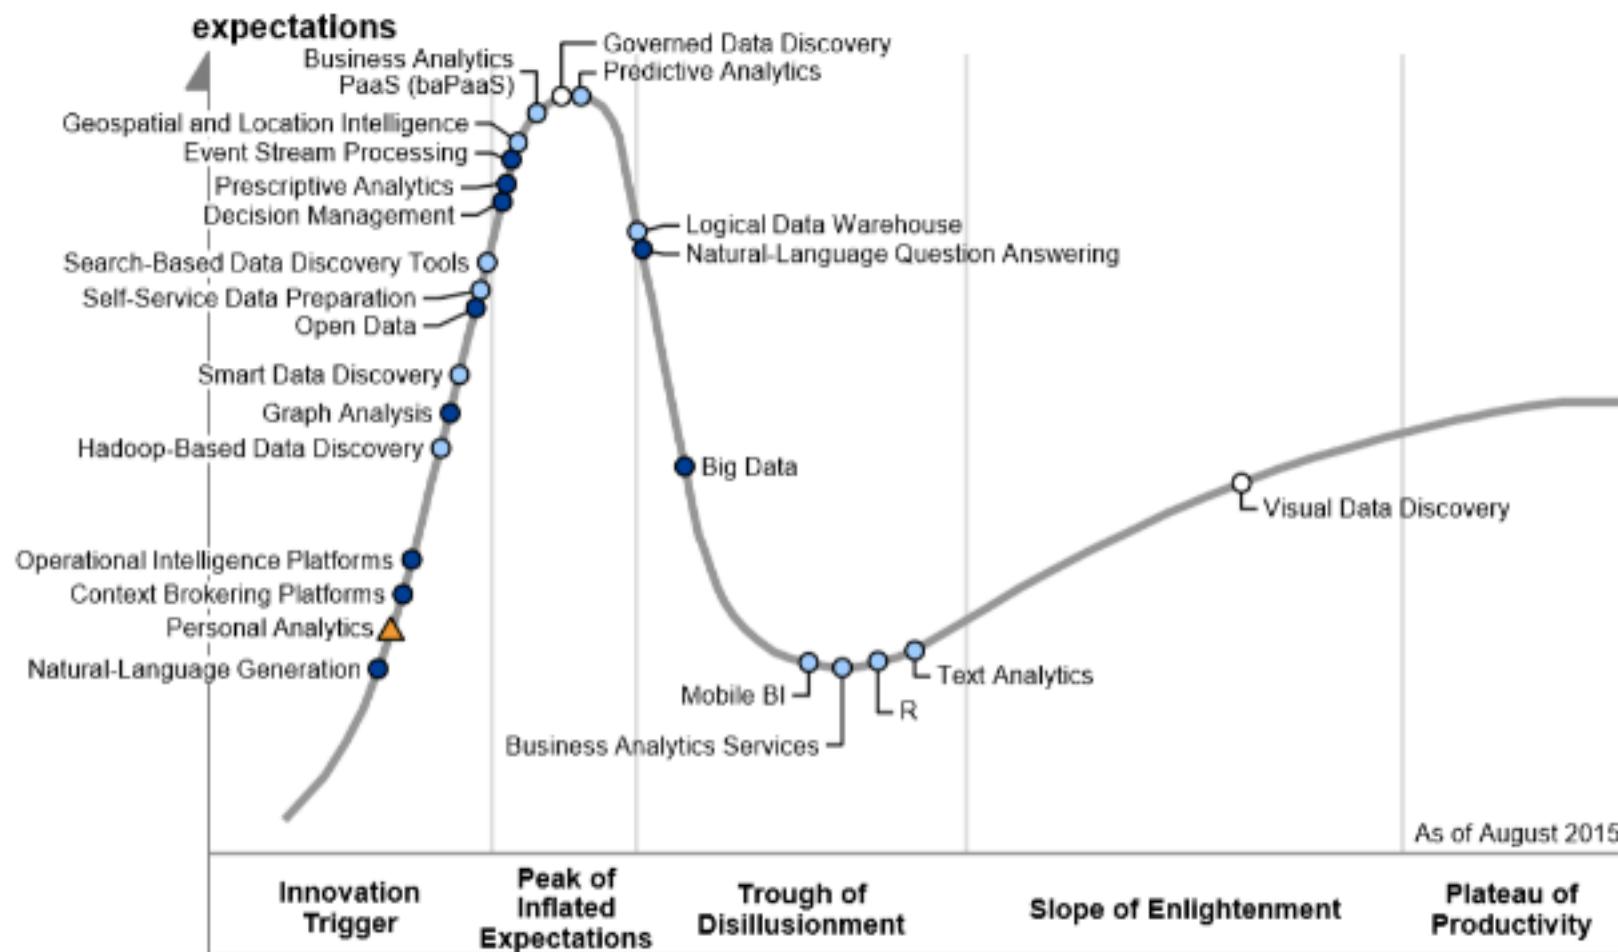
Figure 2. Hype Cycle for Emerging Technologies, 2015



Source: Gartner (July 2015)

BIG DATA OUT OF EMERGING TECHNOLOGIES

Figure 1. Hype Cycle for Business Intelligence and Analytics, 2015



## BIG DATA IN ANALYTICS